

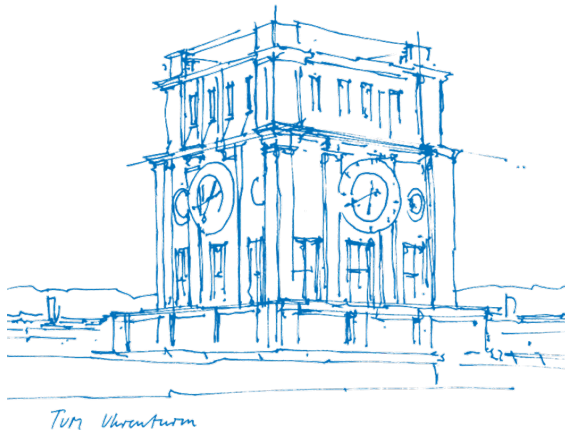
Hist-Tree

An Efficient Indexing Data Structure

Mert Sidal

School of Computation, Information and Technology
Decentralized Information Systems and Data
Management
Technical University of Munich

February 6th, 2025



- 1 Background
- 2 Implementation Details
- 3 Testing & Benchmarking
- 4 Conclusion

Hist-Tree

- **Index** for fast approximate Lookups
- **Basic Assumptions:** Sortedness and Range of Data
- **Idea:** Histogram to partition Data into equal-sized Bins
- Physically organized into **two** Arrays of 32-bit Integers
 1. Inner Nodes **with** Child Pointers
 2. Leaf Nodes

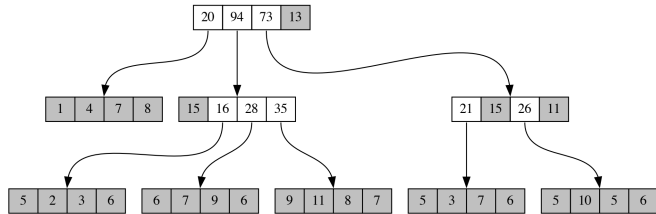


Figure 1 Example Hist-Tree with 200 keys in the range [0,1000)

- 1 Background
- 2 Implementation Details**
- 3 Testing & Benchmarking
- 4 Conclusion

Components

Builder.h:

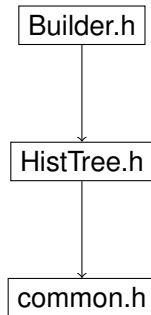
- build()

HistTree.h:

- getSearchBound(key)
- remove(key)
- insert(key)

common.h:

- Utilites: SearchBound struct, Visualizer class and RebuildContext struct



Algorithm 1 HistTree Construction

```
1: create bit vector from sorted keys
2: initial partition into bins
3: while nodes need processing do
4:   count keys in current bins
5:   if bin count < error bound then
6:     stop and create leaf
7:   else
8:     create inner node and split further
9:   end if
10: end while
```

Lookup Algorithm

Algorithm 2 HistTree Lookup

```
1: handle edge cases
2: if inner nodes is empty then
3:   return direct bin lookup in leaf nodes
4: end if
5: while not at leaf do
6:   calculate bin for current level
7:   accumulate counts of previous bins
8:   traverse to next node or return
9:   adjust key and bin width for next level
10: end while
11: return search bound in sorted array
```

Remove Algorithm

Algorithm 3 HistTree Remove

```
1: if key is min/max bound then  
2:   rebuild  
3:   return  
4: end if  
5: check if key exists  
6: reset bit in bit vector  
7: while traversing tree do  
8:   decrement bin count  
9:   if node becomes sparse then  
10:    convert to leaf and cleanup children  
11:   end if  
12: end while
```

Insert Algorithm

Algorithm 4 HistTree Insert

```
1: check if key already exists
2: set bit in bit vector
3: while traversing tree do
4:   increment bin count
5:   if count exceeds error bound then
6:     rebuild
7:     return
8:   end if
9: end while
```

A Collection of Problems & Solutions

Problems

- No clear construction concept in paper
- Limitations of `std::bitset` leading to use of `std::vector<bool>`
- Complexity of keeping the Tree Structure during Updates

Solutions

- Developed custom bit vector approach
- Transitioned to more performant `boost::dynamic_bitset`
- Hybrid Implementation of Updates that trigger a Rebuild **if needed**

Hotspot Analysis: Partition Bit Vector

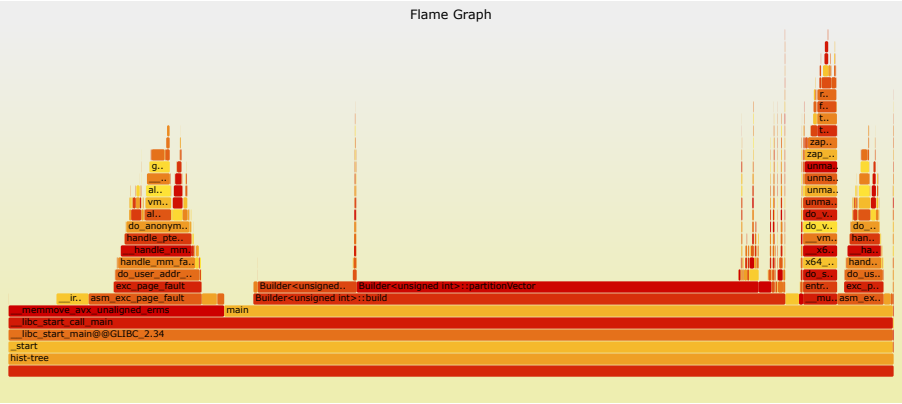


Figure 2 Performance Profiling of Build Function

Bottleneck: partitionVector()

- Repeatedly invoked during Tree Construction
- No Standardized Implementation
- **Base Implementation:** Manual bit-by-bit Copying ($\mathcal{O}(n)$)
- **Optimization Approaches:**
 - memcpy → no support for `boost::dynamic_bitset<>`
 - Bitwise Operations
 - SIMD
 - OpenMP
 - **Final:** Hybrid consisting of SIMD and OpenMP
- **Speedup for large Datasets** approximately 30%

- 1 Background
- 2 Implementation Details
- 3 Testing & Benchmarking**
- 4 Conclusion

Testing Approach

Testing Methodology

- Google Test Framework for Unit-Tests
- **Builder Tests:** Constructor, Bit Vector Computations, Build Mechanism
- **HistTree Tests:** Lookup, Insert, Remove
- **Key Challenge:** No Reference Implementation

Areas for Improvement

- Cover more Edge Cases and Boundary Conditions
- Add Tests for Utility Functions in `common.h`

SOSD Benchmark

Situation

- SOSD: Benchmark Suite for Learned Indexes
- Execution Blocked by Memory Constraints

Why?

- Not enough RAM
- `partitionVector()` and `createBitVector()` allocate vectors/bitsets partly across entire `range_`
- Datasets utilize full Key Range, causing exponential Memory Growth

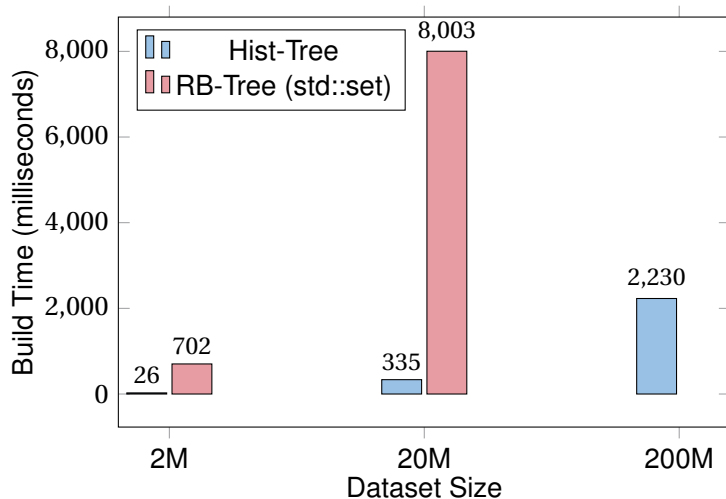
Possible Solutions

- Alternative Build Process without bit vectors
- Bit Vector Compression using Run-Length Encoding (RLE)

Benchmark Approach

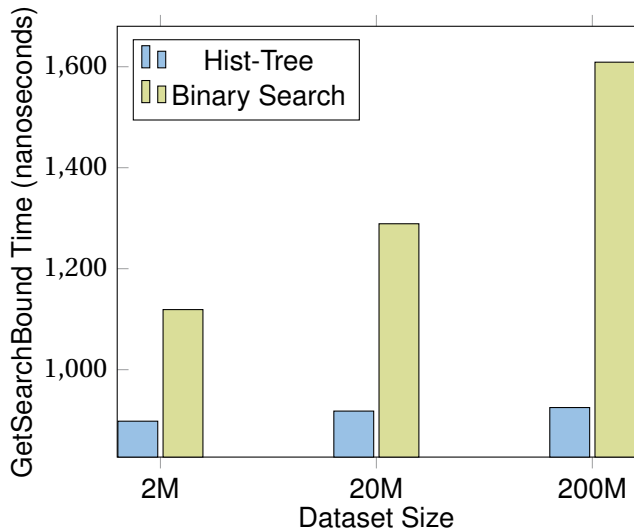
- Benchmarking Framework: Google Benchmark
- Hardware Specifications:
 - CPU: 8 x 2650.12 MHz cores
 - RAM: 8 GB
- Methodology:
 - **Dense** Synthetic Data: sizes 2M–200M
 - Operations Tested:
 - Tree Construction
 - Search Bound Retrieval
 - Insertion & Removal
 - Parameters:
 - Bins: 32–128
 - MaxError: 1024–8192
- **Following Results shown for Bins=64, MaxError=2048**

Benchmark: Construction Time Comparison

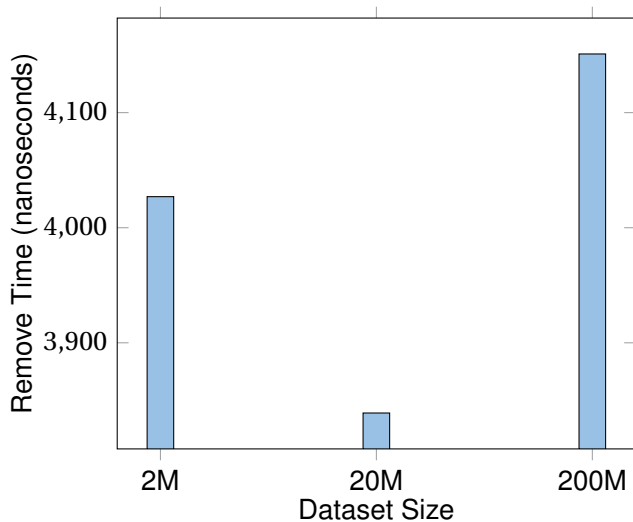


- Hist-Tree Build-Time grows non-linearly
- RB-Tree Construction:
 - 27x slower for 2M
 - 24x slower for 20M
 - Fails at 200M due to memory constraints

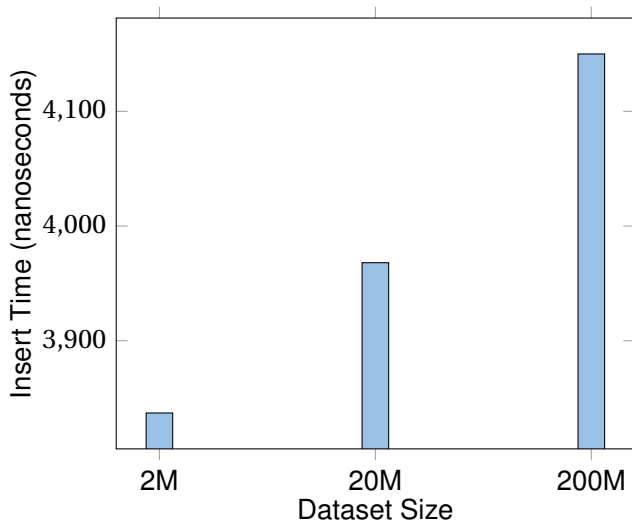
Benchmark: GetSearchBound



Benchmark: Remove

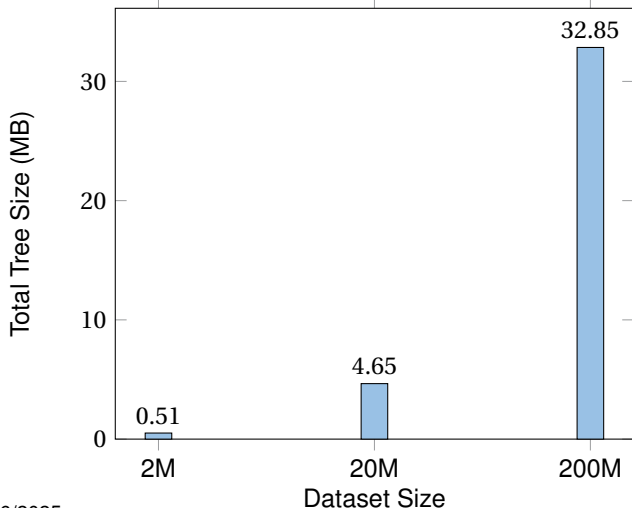


Benchmark: Insert



Benchmark: Memory Footprint of Resulting Hist-Tree

Compression ratio for 200M dataset: 32.846 MB vs. 800 MB raw data size (96% space reduction)



Backup: SOSD Competitors' Memory Footprint

Index / Index Size	XS	S	M	L	XL
	Up to 0.01% of data size	Up to 0.1% of data size	Up to 1% of data size	Up to 10% of data size	No limit
RMI	24.59 KB	24.59 KB	24.59 KB	24.59 KB	24.59 KB
RS	0.17 KB	0.17 KB	0.17 KB	0.17 KB	0.17 KB
PGM	0.37 KB	0.37 KB	0.37 KB	0.37 KB	0.37 KB
BTree	43.07 KB	680.93 KB	5.44 MB	43.54 MB	174.15 MB
FAST		416.77 KB	6.67 MB	6.67 MB	1.71 GB
ALEX		423.1 KB	6.77 MB	54.13 MB	866.1 MB
BinarySearch	0.0 KB	0.0 KB	0.0 KB	0.0 KB	0.0 KB

Figure 3 SOSD Competitors' Memory Footprint (Synthetic Uniform Dense Data (200M))

- 1 Background
- 2 Implementation Details
- 3 Testing & Benchmarking
- 4 Conclusion**

Challenges & Potential Improvements

Current Limitations

- High Memory overhead during Construction
- Costly full Rebuilds during **special** Updates
- Potential Data Fragmentation after multiple Removes

Future Steps

- Optimize Build Process
 - Implement RLE (Run-Length Encoding) approach
 - Rethink the Paper's Build Approach
- Partial Rebuild Strategies
- Defragmentation Techniques
- SOSD: The Revenge

Key Takeaways

- Result has small Memory Footprint
- Near-constant `getSearchBound` (around 900 ns)
- Needs scalable Build